

CSE427s - 2b Lab2

80% (8/10)

- ✓ 1. Recap: the following are the three main components of cloud computing:
- ☐ (A) Compute nodes, Master nodes, and Data nodes
 - ☐ (B) NameNode, JobTracker, and ResourceManager
 - ☒ (C) Storage, Processing, and Resource Management
 - ☐ (D) Batch processing, Stream processing, and Interactive analysis
- ✓ 2. Each *Mapper (Map Task)* in a MapReduce job executes the *map()* function **exactly once**.
- ☐ (A) True
 - ☒ (B) False
- ✗ 3. The *Mapper (Map task)* processes its input data chunk in parallel.
- ☒ (A) True
 - ☐ (B) False
- ✗ 4. One *Reducer (Reduce Task)* gets **all** the values for a particular key.
- ☐ (A) True
 - ☒ (B) False
- ✓ 5. One *Reducer (Reduce Task)* operates exclusively on the values of **one** unique key.
- ☐ (A) True
 - ☒ (B) False
- ✓ 6. In a Hadoop MapReduce execution the *RecordReader* parses the input data into key-value pairs.
- ☒ (A) True
 - ☐ (B) False

- ✓ 7. When using Hadoop MapReduce for Big Data applications these functions are typically implemented by the developer:
- ☐ A readinput()
 - ☒ B map()
 - ☐ C sort()
 - ☐ D group()
 - ☒ E reduce()
 - ☐ F writeoutput()
- ✓ 8. Consider the following data: "*dogs and cats are friends. cats eat cat food or mice.*" The following are results of a word-count MapReduce program:
- ☐ A (cat, 3)
 - ☐ B (cats, 1)
 - ☒ C (dogs, 1)
 - ☒ D (cats, 2)
 - ☒ E (cat, 1)
- ✓ 9. You can terminate a MR job by hitting ctrl+C or closing the terminal window.
- ☐ A True
 - ☒ B False
- ✓ 10. Identify the commonly used Mappers and Reducers:
- ☒ A Filter Mapper
 - ☐ B Filter Reducer
 - ☒ C Sum Reducer
 - ☐ D Average Mapper
 - ☒ E Change Keyspace Mapper