

CSE427s - 3 [P] MR3: Max  
Number of ...

50% (1/2)

✗ 1. The default behaviour in a MapReduce execution is that each Reduce task (Reducer) executes the reduce() function exactly once.

☒ A True

☐ B False

✓ 2. *Maximum parallelism* (executing only one reduce() function per Reduce task) is a good idea.

☐ A Yes

☒ B No

⊘ 3. Give **one argument** why maximum parallelism (executing only one reduce() function per Reduce task) is not good idea.

if one reduce() function per reduce task, there will be so many reduce compute node. in this case, there will be so many small parts of data to be transferred to different nodes. this might take a long time.